

NN-Clustering-WineQuality

November 1, 2020

```
[1]: from sklearn import datasets
from sklearn import preprocessing
from sklearn.model_selection import train_test_split
from sklearn import tree
from sklearn.neural_network import MLPClassifier
from sklearn.ensemble import AdaBoostClassifier
from sklearn import svm
from sklearn.neighbors import KNeighborsClassifier
from sklearn.model_selection import validation_curve
from sklearn.model_selection import GridSearchCV
from sklearn.model_selection import learning_curve
from sklearn.model_selection import cross_validate, cross_val_score
from sklearn.model_selection import KFold
from sklearn.ensemble import RandomForestClassifier as RFC
from sklearn.preprocessing import StandardScaler
from sklearn.decomposition import FastICA as ICA
from sklearn.random_projection import GaussianRandomProjection as RP
from sklearn.decomposition import PCA, FastICA
from sklearn.cluster import KMeans
from sklearn.mixture import GaussianMixture as EM

from sklearn import metrics

import matplotlib.pyplot as plt

import numpy as np
import pandas as pd
import itertools

import time
import timeit
import warnings
warnings.simplefilter('ignore')

[2]: def import_data():
    df_ds1 = pd.read_csv("winequality-white.csv", sep=";").append(pd.
        ↳read_csv("winequality-red.csv", sep=";"), ignore_index=True)
```

```

missing_values = ['?']
df_ds2 = pd.read_csv("breast-cancer-wisconsin.csv", sep=",",na_values =_
↳missing_values)
df_ds2.fillna(method='ffill',inplace=True)
X_ds1 = np.array(df_ds1.values[:,1:-1])
y_ds1 = np.array(df_ds1.values[:, -1])
y_ds1 = (y_ds1<6).astype(int)
X_ds2 = np.array(df_ds2.values[:,1:-1])
y_ds2 = np.array(df_ds2.values[:, -1])
y_ds2 = (y_ds2<3).astype(int)
#     np.where(y_ds2==2,0,y_ds2)
#     np.where(y_ds2==4,1,y_ds2)
return df_ds1, df_ds2, X_ds1, y_ds1, X_ds2, y_ds2

```

```

[3]: def process_data():
    df_Wine, df_BC, X_Wine, Y_Wine, X_BC, Y_BC = import_data()
    #     print(Y_Wine, Y_BC)
    X_Wine = df_Wine.drop('quality', axis=1)
    Y_Wine = df_Wine.quality
    Y_Wine.replace([0,1,2,3,4,5],0,inplace=True)
    Y_Wine.replace([6,7,8,9,10],1,inplace=True)
    #     print(df_BC.head(5))
    X_BC = df_BC.drop('Class', axis=1)
    Y_BC = df_BC.Class
    Y_BC.replace(2,0,inplace=True)
    Y_BC.replace(4,1,inplace=True)
    sc = StandardScaler()
    X_Wine = sc.fit_transform(X_Wine)
    #     print(type(X_Wine))
    X_BC = sc.fit_transform(X_BC)
    return df_Wine, df_BC, X_Wine, Y_Wine, X_BC, Y_BC

```

```

[4]: def train_test(X, y):
    X_train, X_test, y_train, y_test = train_test_split(np.array(X),np.
↳array(y), test_size=0.20)
    sc = preprocessing.StandardScaler()
    X_train = sc.fit_transform(X_train)
    X_test = sc.transform(X_test)
    return X_train, X_test, y_train, y_test

```

0.1 Neural Network Functions

```

[5]: def plot_learning_curve_nn(train_sizes, train_mean, train_std, cv_mean, cv_std,
↳title):
    plt.figure()
    plt.plot(train_sizes, train_mean, '-', color="b", label="Training Score")

```

```

plt.plot(train_sizes, cv_mean, '-', color="r", label="Cross-Validation_
→Score")
plt.title("Learning Curve For Nueral Network: " + title)
plt.xlabel("Training Examples")
plt.ylabel("F1 Score")
plt.legend(loc="best")
plt.show()

```

```

[6]: def plot_times(train_sizes, fit_mean, fit_std, pred_mean, pred_std, title):
plt.figure()
plt.plot(train_sizes, fit_mean, '-', color="b", label="Training Time (s)")
plt.plot(train_sizes, pred_std, '-', color="r", label="Prediction Time (s)")
plt.title("Neural Network model time: " + title)
plt.xlabel("Training Examples")
plt.ylabel("Training Time")
plt.legend(loc="best")
plt.show()

```

```

[7]: def plot_fit_time(n,full_fit,pca_fit,ica_fit,RP_fit,rfc_fit,title):
plt.figure()
plt.plot(n, full_fit, '-', color="k", label="Full Dataset")
plt.plot(n, pca_fit, '-', color="b", label="PCA")
plt.plot(n, ica_fit, '-', color="r", label="ICA")
plt.plot(n, RP_fit, '-', color="g", label="RP")
plt.plot(n, rfc_fit, '-', color="m", label="RFC")
plt.title("Neural Network: training time: " + title)
plt.xlabel("Training Examples")
plt.ylabel("Training Time")
plt.legend(loc="best")
plt.show()

```

```

[8]: def plot_pred_time(n,full_pred, pca_pred, ica_pred, RP_pred, rfc_pred, title):
plt.figure()
plt.plot(n, full_pred, '-', color="k", label="Full Dataset")
plt.plot(n, pca_pred, '-', color="b", label="PCA")
plt.plot(n, ica_pred, '-', color="r", label="ICA")
plt.plot(n, RP_pred, '-', color="g", label="RP")
plt.plot(n, rfc_pred, '-', color="m", label="RFC")
plt.title("Neural Network: prediction times: " + title)
plt.xlabel("Training Examples")
plt.ylabel("Prediction Time")
plt.legend(loc="best")
plt.show()

```

```

[9]: def plot_learn_time(n,full_learn, pca_learn, ica_learn, RP_learn, rfc_learn,
→title):

```

```

plt.figure()
plt.plot(n, full_learn, '-', color="k", label="Full Dataset")
plt.plot(n, pca_learn, '-', color="b", label="PCA")
plt.plot(n, ica_learn, '-', color="r", label="ICA")
plt.plot(n, RP_learn, '-', color="g", label="RP")
plt.plot(n, rfc_learn, '-', color="m", label="RFC")
plt.title("Neural Network: learning rate: " + title)
plt.xlabel("Training Examples")
plt.ylabel("F1 Score")
plt.legend(loc="best")
plt.show()

```

```
[10]: def plot_cv_time(n,full_cv, pca_cv, ica_cv, RP_cv, rfc_cv, title):
```

```

plt.figure()
plt.plot(n, full_cv, '-', color="k", label="Full Dataset")
plt.plot(n, pca_cv, '-', color="b", label="PCA")
plt.plot(n, ica_cv, '-', color="r", label="ICA")
plt.plot(n, RP_cv, '-', color="g", label="RP")
plt.plot(n, rfc_cv, '-', color="m", label="RFC")
plt.title("Neural Network: cross validation rate: " + title)
plt.xlabel("Training Examples")
plt.ylabel("F1 Score")
plt.legend(loc="best")
plt.show()

```

```
[11]: def plot_confusion_matrix_nn(clf_dt_best_params_,X_test, y_test):
    titles_options = [("Confusion Matrix", None),
                       ("Confusion Matrix - Sensitivity and Specificity", 'true')]

    for title, normalize in titles_options:
        disp = metrics.plot_confusion_matrix(clf_dt_best_params_, X_test,
→y_test,
                                           cmap=plt.cm.Blues,
                                           normalize=normalize)

        disp.ax_.set_title(title)

    plt.show()

```

```
[12]: def plot_nn_components(n_components, cv_score, cv_score_em, cv_score_km,
→xlabel, ylabel, title):
    fig1, ax1 = plt.subplots()
    plt.plot(n_components, cv_score, 'o-',label='without label')
    plt.plot(n_components, cv_score_em,'o-', label='with EM label')
    plt.plot(n_components, cv_score_km,'o-', label='with KMeans label')
    plt.xlabel(xlabel)

```

```
plt.ylabel(ylabel)
plt.title(title)
plt.legend(loc="best")
plt.show()
```

```
[13]: def classification_report(clf,X_train, X_test, y_train, y_test):

    start_time = timeit.default_timer()
    clf.fit(X_train, y_train)
    end_time = timeit.default_timer()
    training_time = end_time - start_time

    start_time = timeit.default_timer()
    y_pred = clf.predict(X_test)
    end_time = timeit.default_timer()
    pred_time = end_time - start_time
    auc = metrics.roc_auc_score(y_test, y_pred)

    f1 = metrics.f1_score(y_test,y_pred)
    accuracy = metrics.accuracy_score(y_test,y_pred)
    precision = metrics.precision_score(y_test,y_pred)
    recall = metrics.recall_score(y_test,y_pred)
    cm = metrics.confusion_matrix(y_test,y_pred)

    print("Classification Report")
    print("*****")
    print("Training Time:    "+"{:.5f}".format(training_time))
    print("Prediction Time: "+"{:.5f}\n".format(pred_time))
    print("F1 Score:    "+"{:.2f}".format(f1))
    print("Accuracy:    "+"{:.2f}".format(accuracy)+"      AUC:      "+"{:.2f}".
→format(auc))
    print("Precision: "+"{:.2f}".format(precision)+"      Recall:    "+"{:.2f}".
→format(recall))
    print("*****")
    print(metrics.classification_report(y_test, y_pred))
    print("*****")
    plot_confusion_matrix_nn(clf, X_test, y_test)
```

```
[14]: def evaluate_nn(clf, X, y, title="Insert Title"):
    train_mean = []
    train_std = []
    cv_mean = []
    cv_std = []
    fit_mean = []
    fit_std = []
    pred_mean = []
    pred_std = []
```

```

train_sizes=(np.linspace(.05, 1.0, 50)* len(y) ).astype('int')

for i in train_sizes:
    idx = np.random.randint(X.shape[0], size=i)
    X_subset = X[idx,:]
    y_subset = y[idx]
    scores = cross_validate(clf, X_subset, y_subset, cv=10, scoring='f1',
↪n_jobs=-1, return_train_score=True)

    train_mean.append(np.mean(scores['train_score']))
    train_std.append(np.std(scores['train_score']))
    cv_mean.append(np.mean(scores['test_score']))
    cv_std.append(np.std(scores['test_score']))
    fit_mean.append(np.mean(scores['fit_time']))
    fit_std.append(np.std(scores['fit_time']))
    pred_mean.append(np.mean(scores['score_time']))
    pred_std.append(np.std(scores['score_time']))

train_mean = np.array(train_mean); train_std = np.array(train_std)
cv_mean = np.array(cv_mean); cv_std = np.array(cv_std)
fit_mean = np.array(fit_mean); fit_std = np.array(fit_std)
pred_mean = np.array(pred_mean); pred_std = np.array(pred_std)

plot_learning_curve_nn(train_sizes, train_mean, train_std, cv_mean, cv_std,
↪title)
plot_times(train_sizes, fit_mean, fit_std, pred_mean, pred_std, title)

return train_sizes, train_mean, fit_mean, pred_mean, cv_mean

```

```

[15]: def
↪analyze_nn_components(alg,X,Y,components,classes,df=None,column=None,km_val=None):
↪
    n_classes = classes # 2
    time_pca = []
    n_components = components # range(1,8)
    cv_score = []
    cv_score_em = []
    cv_score_km = []
    X_transformed = None
    clust = None
    clust_km = None
    clust_em = None
    if km_val == None:
        km_components =
↪KMeans(n_clusters=n_classes,n_init=10,random_state=100,n_jobs=-1).fit(X)
    else:

```

```

        km_components = km_val
        km_labels = km_components.labels_
#         print('km_labels', km_labels)

#         nodes_hidden_layer = int((comp + n_classes)/2)
        nodes_hidden_layer = (10,5,5)
        #neural network learner
#         mlp = _
        ↪MLPClassifier(hidden_layer_sizes=(nodes_hidden_layer,), max_iter=5000, random_state=18)
        mlp = _
        ↪MLPClassifier(hidden_layer_sizes=nodes_hidden_layer, learning_rate_init=0.
        ↪01, max_iter=5000, random_state=18)

        for comp in n_components:

#             #         nodes_hidden_layer = int((comp + n_classes)/2)
#             nodes_hidden_layer = (10,5,5)
#             #neural network learner
#             #         mlp = _
            ↪MLPClassifier(hidden_layer_sizes=(nodes_hidden_layer,), max_iter=5000, random_state=18)
#             mlp = _
            ↪MLPClassifier(hidden_layer_sizes=nodes_hidden_layer, learning_rate_init=0.
            ↪01, max_iter=5000, random_state=18)

            if alg == 'PCA':
                X_transformed = pca_cluster(X, comp, 18)
            elif alg == 'ICA':
                X_transformed = ica_cluster(X, comp, 18)
            elif alg == 'RP':
                X_transformed = RP_cluster(X, comp, 18)
            elif alg == 'RFC':
                imp_rf, top_cols_rf = run_RFC(X, Y, df.drop(column, axis=1), 500, 5, -1)
                X_transformed = rf_cluster(df, column, top_cols_rf)

            cv_score.append(np.mean(cross_val_score(mlp, X_transformed, Y, cv = 3)))

#         print('km: ', km, 'km_labels: ', km_labels)
        em = _
        ↪EM(n_components=comp, covariance_type='full', n_init=1, warm_start=True, random_state=18).
        ↪fit(X)

        em_labels = em.predict(X)
#         print('em: ', em, 'em_labels: ', em_labels)
#         print('km_labels', km_labels)
        clust = get_clusters(X_transformed, km_labels, em_labels)

```

```

    clust_km = clust[:,0:1]
    clust_em = clust[:,1:2]

    X_transformed_em = np.concatenate((X_transformed, clust_em), axis=1)

    cv_score_em.append(np.mean(cross_val_score(mlp, X_transformed_em, Y, cv_
    ↪= 10)))
    #     cv_score_em_pca.append(np.mean(cross_val_score(mlp,
    ↪X_transformed_em_pca, y_ds2)))

    X_transformed_km = np.concatenate((X_transformed, clust_km), axis=1)

    cv_score_km.append(np.mean(cross_val_score(mlp, X_transformed_km, Y, cv_
    ↪= 10)))
    #     cv_score_km_pca.append(np.mean(cross_val_score(mlp,
    ↪X_transformed_km_pca, y_ds2)))
    return n_components, cv_score, cv_score_em, cv_score_km

```

```

[16]: def get_clusters(X,km_labels,em_labels):
    __X,__km_labels,__em_labels = X,km_labels,em_labels
    #     print('X: ',__X,'km_labels: ', __km_labels,'em_labels: ', __em_labels)
    df = pd.DataFrame(__X)
    df['KM Cluster'] = __km_labels
    df['EM Cluster'] = __em_labels
    col_1hot = ['KM Cluster', 'EM Cluster']
    df_1hot = df[col_1hot]
    df_1hot = pd.get_dummies(df_1hot).astype('category')
    df_others = df.drop(col_1hot,axis=1)
    df = pd.concat([df_others,df_1hot],axis=1)
    #     df = pd.concat([df_1hot],axis=1)
    X_cluster = np.array(df.values, dtype='int64')

    return X_cluster

```

```

[17]: def run_RFC(X,y,df_original,n_estimators,random_state,n_jobs):
    rfc = RFC(n_estimators=n_estimators,min_samples_leaf=round(len(X)*.
    ↪01),random_state=random_state,n_jobs=n_jobs)
    imp = rfc.fit(X,y).feature_importances_
    imp = pd.DataFrame(imp,columns=['Feature Importance'],index=df_original.
    ↪columns)
    imp.sort_values(by=['Feature Importance'],inplace=True,ascending=False)
    imp['Cum Sum'] = imp['Feature Importance'].cumsum()

```



```

imp = imp[imp['Cum Sum']<=0.87]
top_cols = imp.index.tolist()
return imp, top_cols

```

```

[18]: def pca_cluster(X,n_components,random_state):
        pca_cluster = PCA(n_components=n_components,random_state=random_state).
        ↪fit_transform(X)
        return pca_cluster

```

```

[19]: def ica_cluster(X,n_components,random_state):
        ica_cluster = FastICA(n_components = n_components,whiten=True,
        ↪random_state= random_state).fit_transform(X)
        return ica_cluster

```

```

[20]: def RP_cluster(X,n_components,random_state):
        RP_cluster = RP(n_components = n_components, random_state= random_state).
        ↪fit_transform(X)
        return RP_cluster

```

```

[21]: def rf_cluster(df_ds, column, top_cols):
        rf_cluster = df_ds.drop(column, axis=1)[top_cols]
        return rf_cluster

```

1 Load Datasets

```

[22]: # df_Wine, df_BC, X_Wine, Y_Wine, X_BC, Y_BC = process_data()

df_BC, df_Wine, X_BC, Y_BC, X_Wine, Y_Wine = process_data()

```

```

[23]: print(X_BC)

```

```

[[-0.16608919 -0.42318303  0.28468605 ... -1.35904886 -0.54617826
  -1.41855821]
 [-0.70607349 -0.24094936  0.14704613 ...  0.50691489 -0.27735097
  -0.83161516]
 [ 0.68245757 -0.36243847  0.55996589 ...  0.25811972 -0.61338508
  -0.32852111]
 ...
 [-0.70607349  1.03468634 -1.29817304 ...  1.25330039  1.47002637
  0.42611996]
 [-1.01463595  1.85473786 -1.366993 ...  2.18628226  1.20119908
  -0.2446721 ]
 [-0.93749534 -0.1802048  1.04170561 ...  1.06670401  0.86516498
  0.42611996]]

```

```

[24]: print((np.linspace(.05, 1.0, 20)* (len(Y_BC) ) ).astype('int') )

```

```
[ 324  649  974 1299 1624 1949 2273 2598 2923 3248 3573 3898 4223 4547
 4872 5197 5522 5847 6172 6497]
```

```
[ ]:
```

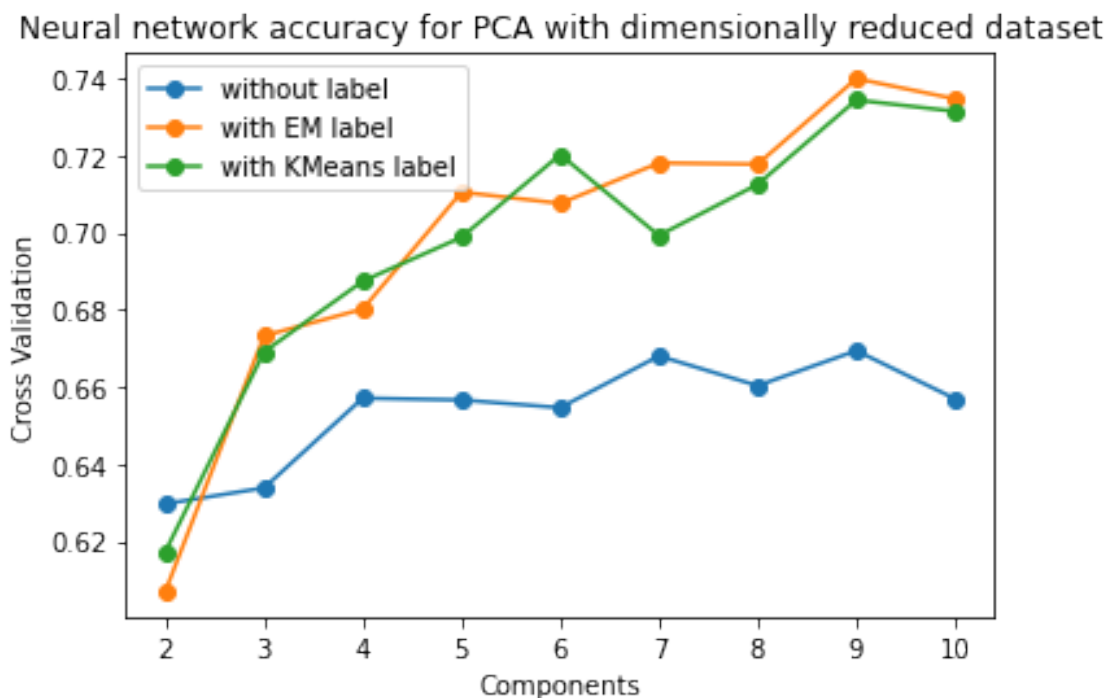
2 Neural Network - Component Analysis

```
[ ]:
```

PCA

```
[25]: km_components = KMeans(n_clusters=4,n_init=10,random_state=18,n_jobs=-1).
      ↪fit(X_BC)
      n_components_pca, cv_score_pca, cv_score_em_pca, cv_score_km_pca =
      ↪analyze_nn_components('PCA',X_BC,Y_BC,range(2,11),4,km_val=km_components)
```

```
[26]: plot_nn_components(n_components_pca, cv_score_pca, cv_score_em_pca,
      ↪cv_score_km_pca, 'Components', 'Cross Validation', 'Neural network accuracy'
      ↪for PCA with dimensionally reduced dataset')
```




ICA

```
[27]: km_components = KMeans(n_clusters=4,n_init=10,random_state=18,n_jobs=-1).
      ↪fit(X_BC)
```

```
n_components_ica, cv_score_ica, cv_score_em_ica, cv_score_km_ica =  
→analyze_nn_components('ICA',X_BC,Y_BC,range(2,11),4,km_val=km_components)
```

```
[28]: plot_nn_components(n_components_ica, cv_score_ica, cv_score_em_ica,  
→cv_score_km_ica, 'Components', 'Cross Validation', 'Neural network accuracy',  
→for ICA with dimensionally reduced dataset')
```




Randomize Projection

```
[29]: km_components = KMeans(n_clusters=4,n_init=10,random_state=100,n_jobs=-1).  
→fit(X_BC)  
n_components_RP, cv_score_RP, cv_score_em_RP, cv_score_km_RP =  
→analyze_nn_components('RP',X_BC,Y_BC,range(2,11),4,km_val=km_components)
```

```
[30]: plot_nn_components(n_components_RP, cv_score_RP, cv_score_em_RP,  
→cv_score_km_RP, 'Components', 'Cross Validation', 'Neural network accuracy',  
→for Randomize Projection with dimensionally reduced dataset')
```

Neural network accuracy for Randomize Projection with dimensionally reduced dataset



Random Forest Classifier

[31]: `km_components = KMeans(n_clusters=4,n_init=10,random_state=100,n_jobs=-1).`

`→fit(X_BC)`

`n_components_rfc, cv_score_rfc, cv_score_em_rfc, cv_score_km_rfc =`

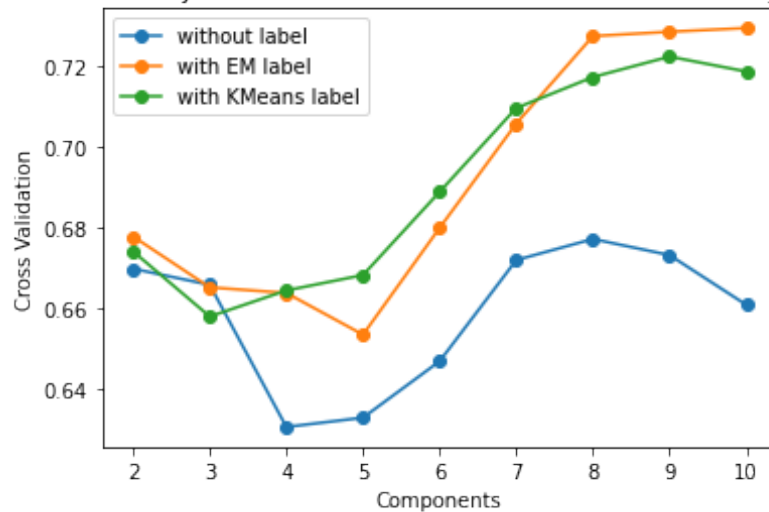
`→analyze_nn_components('RFC',X_BC,Y_BC,range(2,11),4,df=df_BC,column='quality',km_val=km_com`

[32]: `plot_nn_components(n_components_RP, cv_score_RP, cv_score_em_RP,`

`→cv_score_km_RP, 'Components', 'Cross Validation', 'Neural network accuracy,`

`→for Random Forest Classifier with dimensionally reduced dataset')`

Neural network accuracy for Random Forest Classifier with dimensionally reduced dataset



[]:

3 Training Neural Network on Projected Data

```
[33]: # imp_rf_BC, top_cols_rf_BC = run_RFC(X_BC, Y_BC, df_BC.drop('Class',  
→axis=1), 500, 5, -1)  
imp_rf_BC, top_cols_rf_BC = run_RFC(X_BC, Y_BC, df_BC.drop('quality',  
→axis=1), 500, 5, -1)
```

```
[34]: # X_pca_cluster_BC = pca_cluster(X_BC, 6, 5)  
# X_ica_cluster_BC = ica_cluster(X_BC, 12, 5)  
# X_RP_cluster_BC = RP_cluster(X_BC, 8, 5)  
# # X_rf_cluster_BC = rf_cluster(df_BC, 'Class', top_cols_rf_BC)  
# X_rf_cluster_BC = rf_cluster(df_BC, 'quality', top_cols_rf_BC)
```

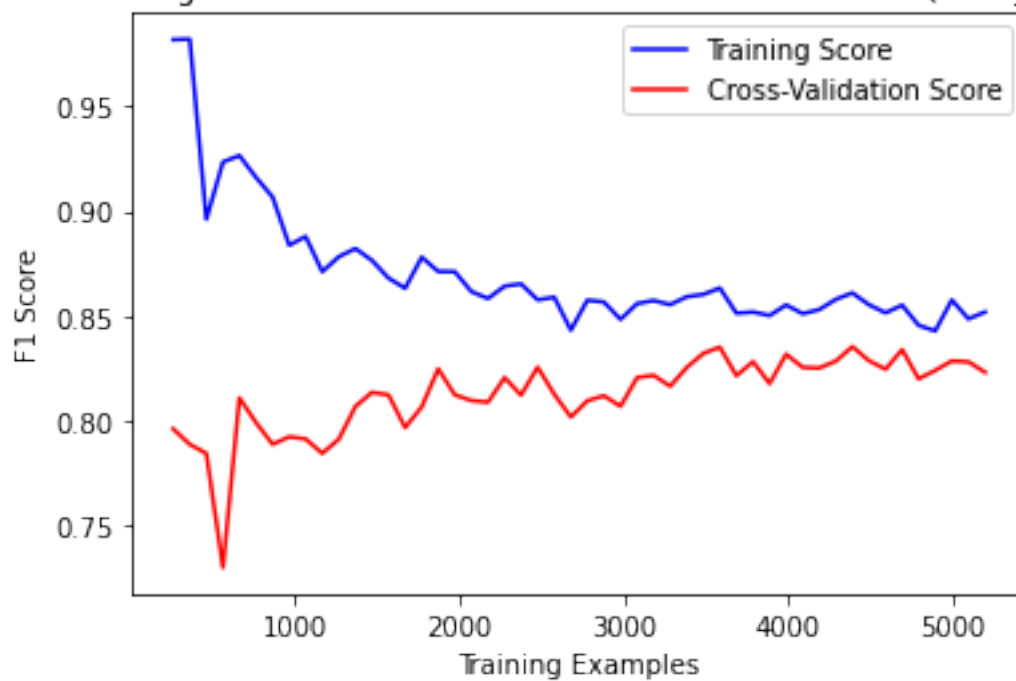
```
[35]: X_pca_cluster_BC = pca_cluster(X_BC, 4, 18)  
X_ica_cluster_BC = ica_cluster(X_BC, 4, 18)  
X_RP_cluster_BC = RP_cluster(X_BC, 4, 18)  
# X_rf_cluster_BC = rf_cluster(df_BC, 'Class', top_cols_rf_BC)  
X_rf_cluster_BC = rf_cluster(df_BC, 'quality', top_cols_rf_BC)
```

```
[36]: # X_pca_cluster_BC = pca_cluster(X_BC, 8, 5)  
# X_ica_cluster_BC = ica_cluster(X_BC, 5, 5)  
# X_RP_cluster_BC = RP_cluster(X_BC, 5, 5)  
# # X_rf_cluster_BC = rf_cluster(df_BC, 'Class', top_cols_rf_BC)  
# X_rf_cluster_BC = rf_cluster(df_BC, 'quality', top_cols_rf_BC)
```

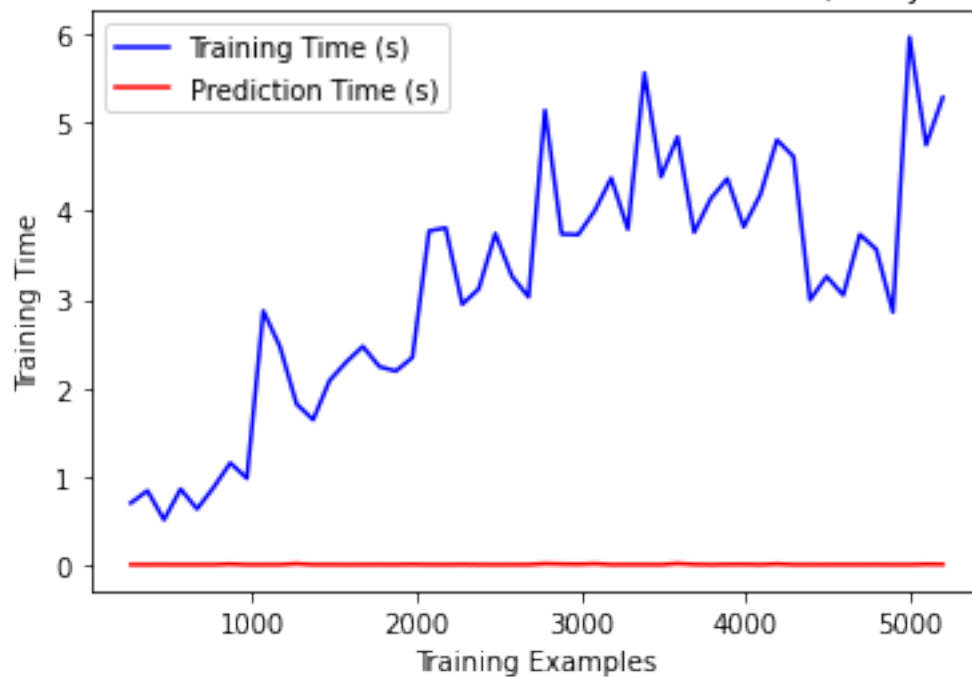
Full Dataset

```
[37]: X_train_nn, X_test_nn, y_train_nn, y_test_nn = train_test(X_BC, Y_BC)  
full_est = MLPClassifier(hidden_layer_sizes=(10, 5, 5), solver='adam',  
→activation='relu', learning_rate_init=0.01, max_iter=10000, random_state=18)  
train_samp_full, NN_train_score_full, NN_fit_time_full, NN_pred_time_full,  
→NN_cv_score_full = evaluate_nn(full_est, X_train_nn,  
→y_train_nn, title="Neural Net Wine Quality: Full")  
classification_report(full_est, X_train_nn, X_test_nn, y_train_nn, y_test_nn)
```

Learning Curve For Nueral Network: Neural Net Wine Quality: Full



Neural Network model time: Neural Net Wine Quality: Full



Classification Report

Training Time: 3.27753


Prediction Time: 0.00087


F1 Score: 0.82

Accuracy: 0.78 AUC: 0.77

Precision: 0.83 Recall: 0.82

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.70 | 0.72 | 0.71 | 490 |
| 1 | 0.83 | 0.82 | 0.82 | 810 |
| accuracy | | | 0.78 | 1300 |
| macro avg | 0.77 | 0.77 | 0.77 | 1300 |
| weighted avg | 0.78 | 0.78 | 0.78 | 1300 |







PCA

```
[38]: X_train_nn, X_test_nn, y_train_nn, y_test_nn = train_test(X_pca_cluster_BC,
    ↪Y_BC)
pca_est = MLPClassifier(hidden_layer_sizes=(10,5,5), solver='adam',
    ↪activation='relu', learning_rate_init=0.01, max_iter=10000, random_state=18)
train_samp_pca, NN_train_score_pca, NN_fit_time_pca, NN_pred_time_pca,
    ↪NN_cv_score_pca = evaluate_nn(pca_est, X_train_nn, y_train_nn, title="Neural
    ↪Net Wine Quality: PCA")
classification_report(pca_est, X_train_nn, X_test_nn, y_train_nn, y_test_nn)
```


Learning Curve For Nueral Network: Neural Net Wine Quality: PCA



Neural Network model time: Neural Net Wine Quality: PCA



Classification Report

Training Time: 1.27316


Prediction Time: 0.00241


F1 Score: 0.77

Accuracy: 0.71 AUC: 0.70

Precision: 0.79 Recall: 0.76

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.59 | 0.64 | 0.61 | 459 |
| 1 | 0.79 | 0.76 | 0.77 | 841 |
| accuracy | | | 0.71 | 1300 |
| macro avg | 0.69 | 0.70 | 0.69 | 1300 |
| weighted avg | 0.72 | 0.71 | 0.72 | 1300 |

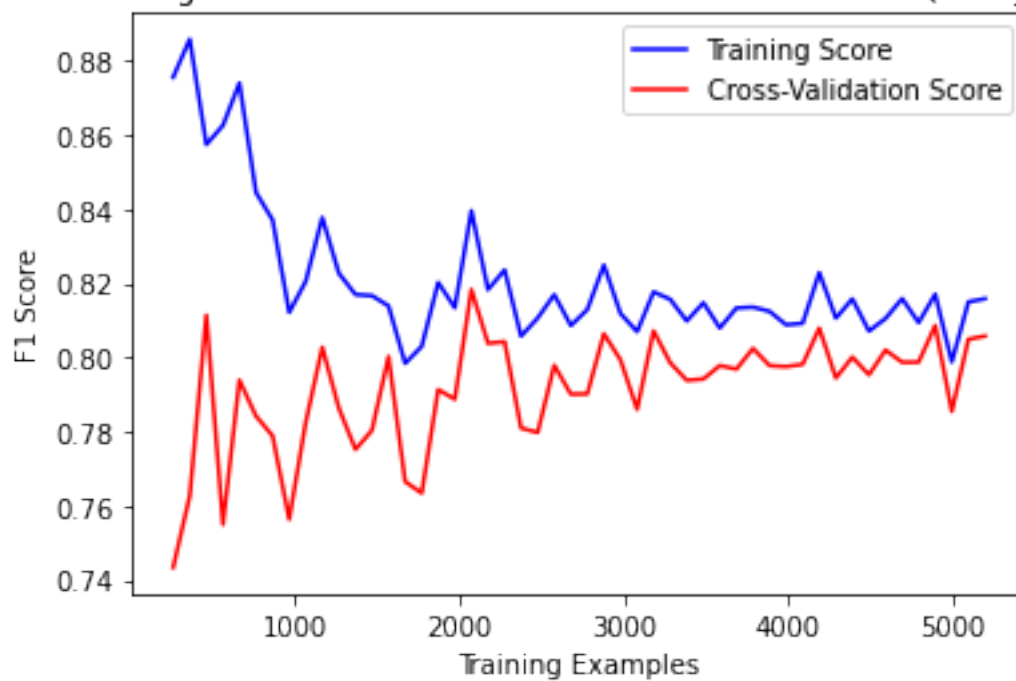




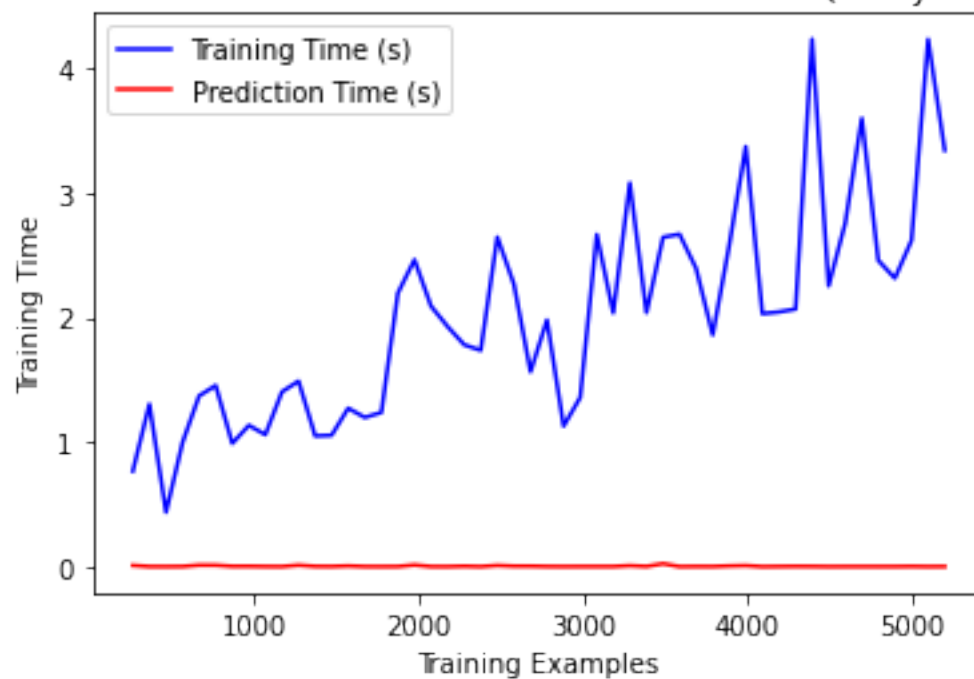
ICA

```
[39]: X_train_nn, X_test_nn, y_train_nn, y_test_nn = train_test(X_ica_cluster_BC,
    ↪Y_BC)
ica_est = MLPClassifier(hidden_layer_sizes=(10,5,5), solver='adam',
    ↪activation='relu', learning_rate_init=0.01, max_iter=10000, random_state=18)
train_samp_ica, NN_train_score_ica, NN_fit_time_ica, NN_pred_time_ica,
    ↪NN_cv_score_ica = evaluate_nn(ica_est, X_train_nn, y_train_nn, title="Neural
    ↪Net Wine Quality: ICA")
classification_report(ica_est, X_train_nn, X_test_nn, y_train_nn, y_test_nn )
```

Learning Curve For Nueral Network: Neural Net Wine Quality: ICA



Neural Network model time: Neural Net Wine Quality: ICA



Classification Report

Training Time: 1.36220


Prediction Time: 0.00084


F1 Score: 0.79

Accuracy: 0.73 AUC: 0.70

Precision: 0.76 Recall: 0.83

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.66 | 0.57 | 0.61 | 490 |
| 1 | 0.76 | 0.83 | 0.79 | 810 |
| accuracy | | | 0.73 | 1300 |
| macro avg | 0.71 | 0.70 | 0.70 | 1300 |
| weighted avg | 0.72 | 0.73 | 0.72 | 1300 |

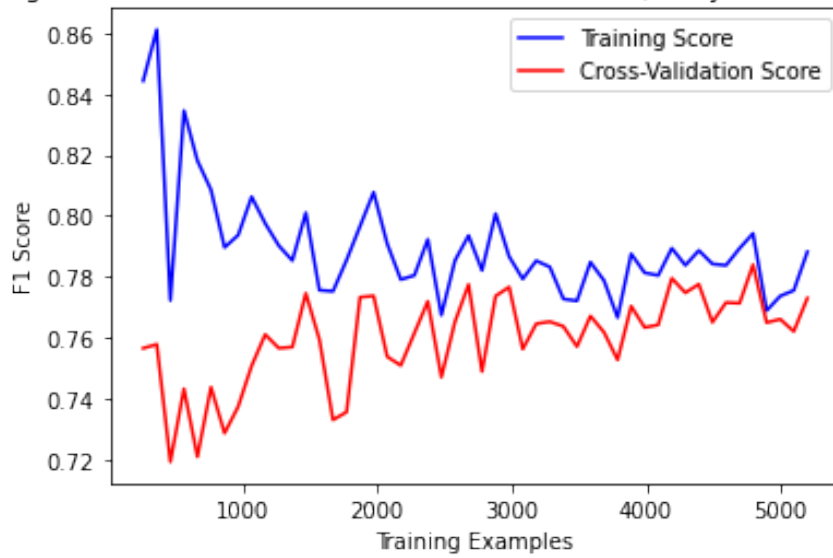





Randomize Projection

```
[40]: X_train_nn, X_test_nn, y_train_nn, y_test_nn = train_test(X_RP_cluster_BC, Y_BC)
      RP_est = MLPClassifier(hidden_layer_sizes=(10,5,5), solver='adam',
      ↪activation='relu', learning_rate_init=0.01, max_iter=10000, random_state=18)
      train_samp_RP, NN_train_score_RP, NN_fit_time_RP, NN_pred_time_RP,
      ↪NN_cv_score_RP = evaluate_nn(RP_est, X_train_nn, y_train_nn, title="Neural
      ↪Net Wine Quality: Randomize Projection")
      classification_report(RP_est, X_train_nn, X_test_nn, y_train_nn, y_test_nn )
```

Learning Curve For Nueral Network: Neural Net Wine Quality: Randomize Projection



Neural Network model time: Neural Net Wine Quality: Randomize Projection



Classification Report

Training Time: 2.45994

Prediction Time: 0.00095

F1 Score: 0.76

```


Accuracy: 0.70      AUC:      0.67
Precision: 0.75     Recall:    0.77
*****
      precision    recall  f1-score   support

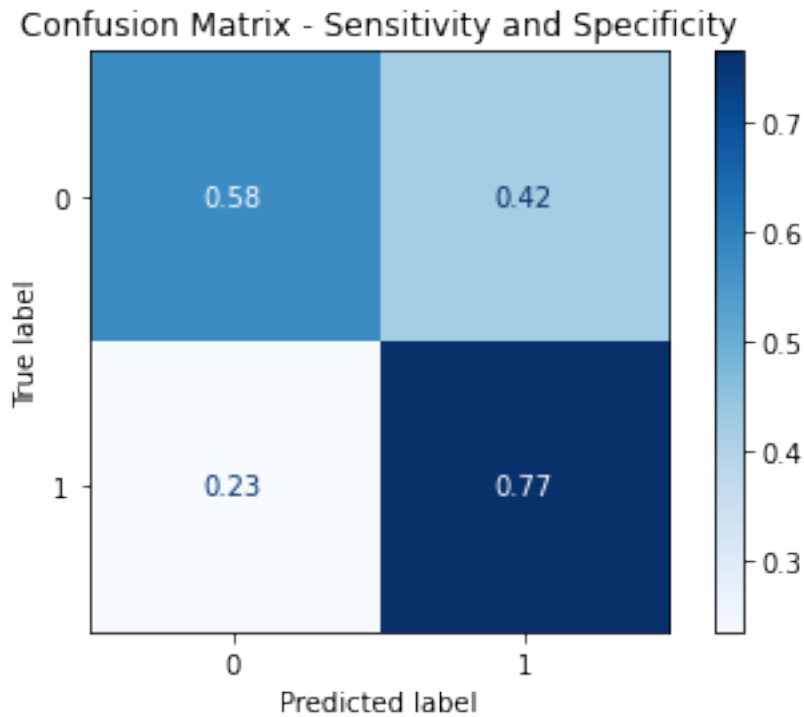
     0       0.60      0.58      0.59       487
     1       0.75      0.77      0.76       813

 accuracy          0.70       1300
 macro avg       0.67      0.67      0.67       1300
 weighted avg    0.69      0.70      0.69       1300

*****

```

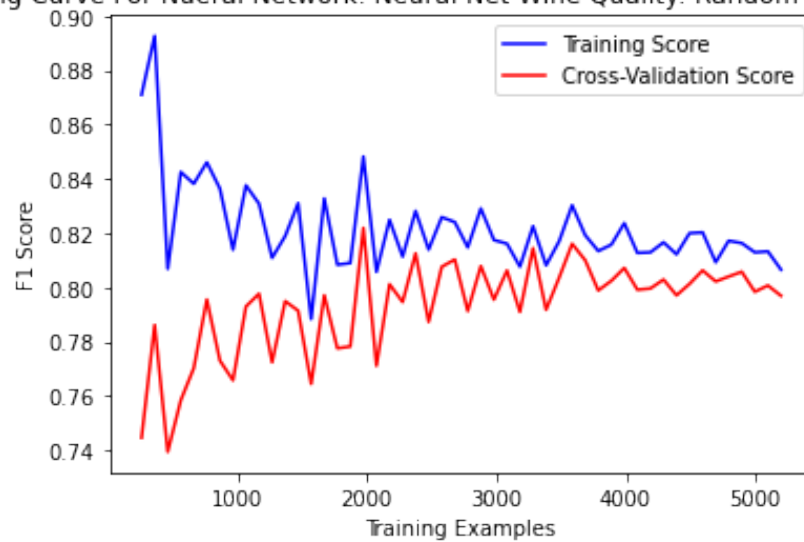





Random Forest Classifier

```
[41]: X_train_nn, X_test_nn, y_train_nn, y_test_nn = train_test(X_ica_cluster_BC,
    ↪Y_BC)
rfc_est = MLPClassifier(hidden_layer_sizes=(10,5,5), solver='adam',
    ↪activation='relu', learning_rate_init=0.01, max_iter=10000, random_state=18)
train_samp_rfc, NN_train_score_rfc, NN_fit_time_rfc, NN_pred_time_rfc,
    ↪NN_cv_score_rfc = evaluate_nn(rfc_est, X_train_nn, y_train_nn, title="Neural
    ↪Net Wine Quality: Random Forest Classifier")
classification_report(rfc_est, X_train_nn, X_test_nn, y_train_nn, y_test_nn)
```

Learning Curve For Nueral Network: Neural Net Wine Quality: Random Forest Classifier



Neural Network model time: Neural Net Wine Quality: Random Forest Classifier



Classification Report

Training Time: 4.80027


Prediction Time: 0.00150


F1 Score: 0.78

Accuracy: 0.70 AUC: 0.67

Precision: 0.74 Recall: 0.81

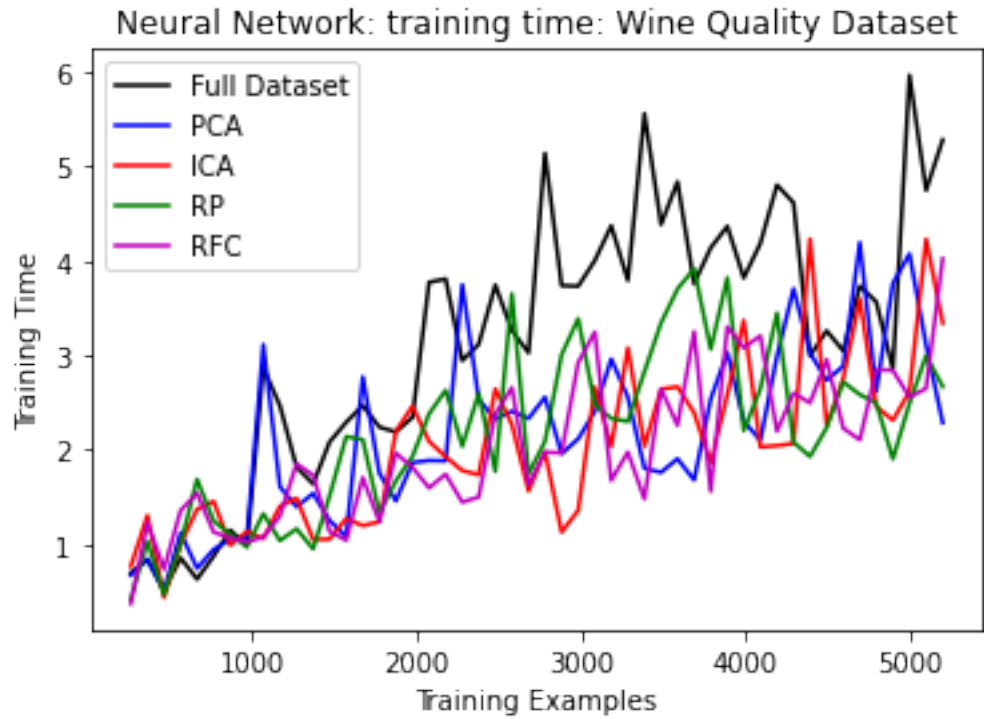
| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.62 | 0.52 | 0.57 | 482 |
| 1 | 0.74 | 0.81 | 0.78 | 818 |
| accuracy | | | 0.70 | 1300 |
| macro avg | 0.68 | 0.67 | 0.67 | 1300 |
| weighted avg | 0.70 | 0.70 | 0.70 | 1300 |



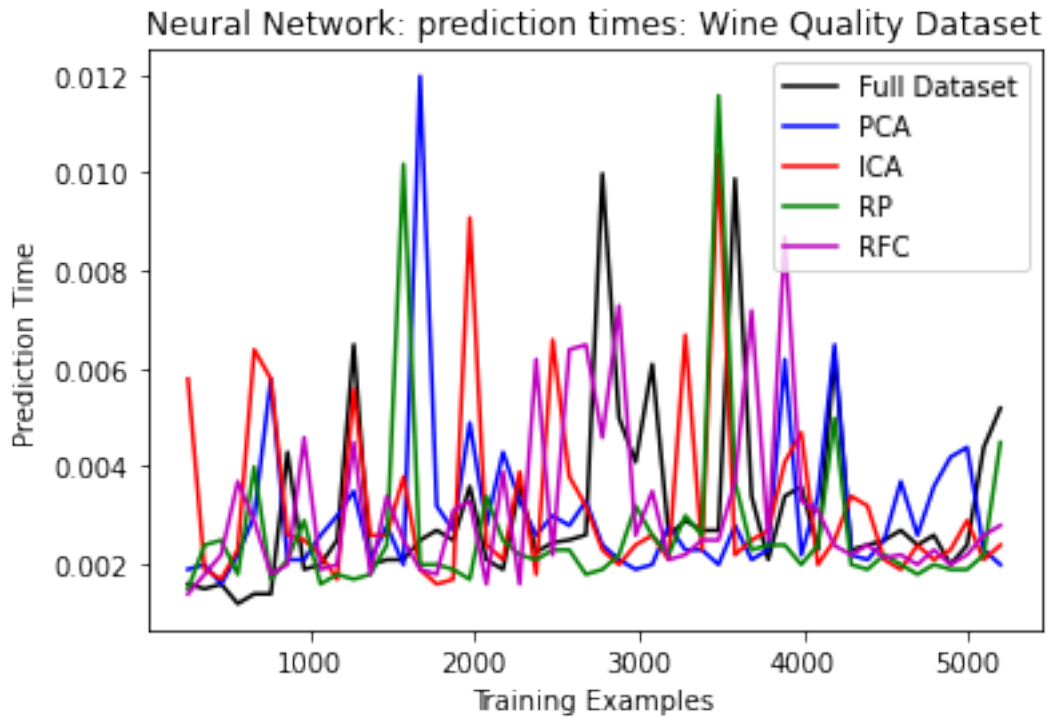


Comparison plots

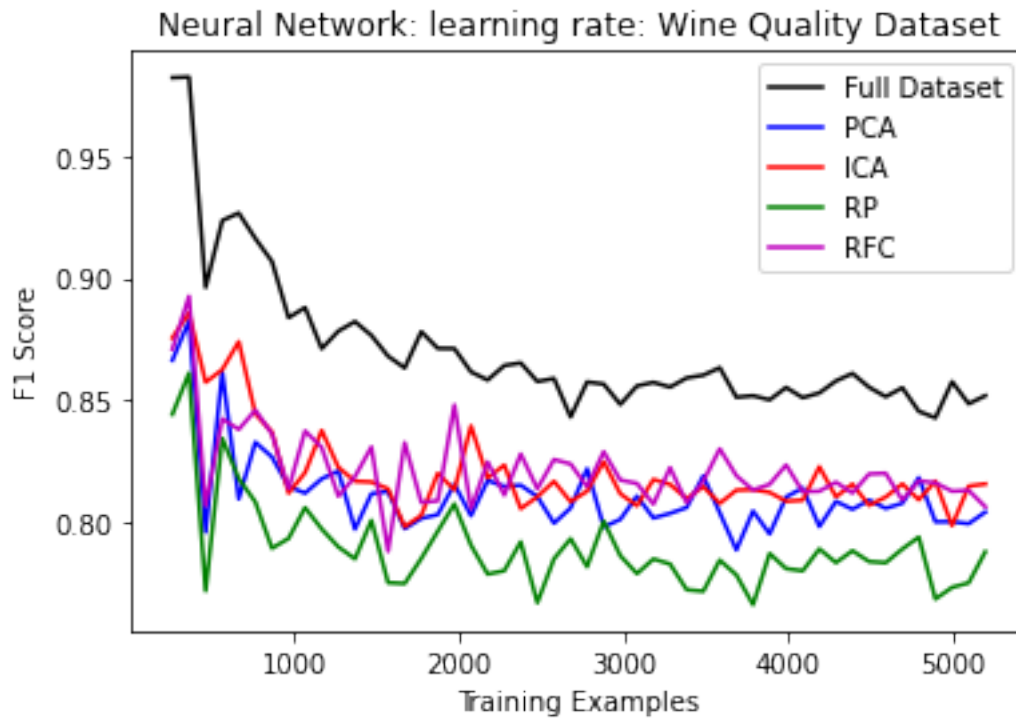
```
[42]: plot_fit_time(train_samp_full, NN_fit_time_full, NN_fit_time_pca, ↵  
    ↪ NN_fit_time_ica, NN_fit_time_RP, NN_fit_time_rfc, 'Wine Quality Dataset')
```




```
[43]: plot_pred_time(train_samp_full, NN_pred_time_full, NN_pred_time_pca,
↪NN_pred_time_ica, NN_pred_time_RP, NN_pred_time_rfc, 'Wine Quality Dataset')
```



```
[44]: plot_learn_time(train_samp_full, NN_train_score_full, NN_train_score_pca,
    ↪ NN_train_score_ica, NN_train_score_RP, NN_train_score_rfc, 'Wine Quality_
    ↪ Dataset')
```



```
[45]: plot_cv_time(train_samp_full, NN_cv_score_full, NN_cv_score_pca,
    ↪ NN_cv_score_ica, NN_cv_score_RP, NN_cv_score_rfc, 'Wine Quality Dataset')
```



4 Training Neural Network on Projected Data with Cluster Labels


```
[46]: # km = KMeans(n_clusters=9,n_init=10,random_state=100,n_jobs=-1).fit(X_BC)
# km_labels = km.labels_
# em =
→EM(n_components=24,covariance_type='diag',n_init=1,warm_start=True,random_state=100).
→fit(X_BC)
# em_labels = em.predict(X_BC)


km = KMeans(n_clusters=5,n_init=10,random_state=18,n_jobs=-1).fit(X_BC)
km_labels = km.labels_
em =
→EM(n_components=7,covariance_type='diag',n_init=1,warm_start=True,random_state=18).
→fit(X_BC)
em_labels = em.predict(X_BC)

clust_full = get_clusters(X_BC,km_labels,em_labels)
clust_pca = get_clusters(X_pca_cluster_BC,km_labels,em_labels)
clust_ica = get_clusters(X_ica_cluster_BC,km_labels,em_labels)
clust_RP = get_clusters(X_RP_cluster_BC,km_labels,em_labels)
clust_rfc = get_clusters(X_rf_cluster_BC,km_labels,em_labels)
```

Full Cluster

```
[47]: X_train_nn, X_test_nn, y_train_nn, y_test_nn = train_test(clust_full, Y_BC)
full_est_lbl = MLPClassifier(hidden_layer_sizes=(10,5,5), solver='adam',
    ↪activation='relu', learning_rate_init=0.01, max_iter=10000, random_state=18)
train_samp_full_lbl, NN_train_score_full_lbl, NN_fit_time_full_lbl,
    ↪NN_pred_time_full_lbl, NN_cv_score_full_lbl = evaluate_nn(full_est_lbl,
    ↪X_train_nn, y_train_nn, title="Neural Net Wine: Full")
classification_report(full_est_lbl, X_train_nn, X_test_nn, y_train_nn,
    ↪y_test_nn)
```





Classification Report

Training Time: 1.69500



Prediction Time: 0.00116

F1 Score: 0.79

Accuracy: 0.71 AUC: 0.66


Precision: 0.74 Recall: 0.84


| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.63 | 0.49 | 0.55 | 470 |
| 1 | 0.74 | 0.84 | 0.79 | 830 |
| accuracy | | | 0.71 | 1300 |
| macro avg | 0.69 | 0.66 | 0.67 | 1300 |
| weighted avg | 0.70 | 0.71 | 0.70 | 1300 |



PCA

```
[48]: X_train_nn, X_test_nn, y_train_nn, y_test_nn = train_test(clust_pca, Y_BC)
pca_est_lbl = MLPClassifier(hidden_layer_sizes=(10,5,5), solver='adam',
    ↪activation='relu', learning_rate_init=0.01, max_iter=10000, random_state=18)
train_samp_pca_lbl, NN_train_score_pca_lbl, NN_fit_time_pca_lbl,
    ↪NN_pred_time_pca_lbl, NN_cv_score_pca_lbl = evaluate_nn(pca_est_lbl,
    ↪X_train_nn, y_train_nn, title="Neural Net Wine Quality: PCA")
classification_report(pca_est_lbl, X_train_nn, X_test_nn, y_train_nn, y_test_nn)
```





Classification Report

Training Time: 1.78400



Prediction Time: 0.00078

F1 Score: 0.78

Accuracy: 0.70 AUC: 0.66


Precision: 0.74 Recall: 0.82


| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.62 | 0.51 | 0.56 | 480 |
| 1 | 0.74 | 0.82 | 0.78 | 820 |
| accuracy | | | 0.70 | 1300 |
| macro avg | 0.68 | 0.66 | 0.67 | 1300 |
| weighted avg | 0.69 | 0.70 | 0.69 | 1300 |



ICA

```
[49]: X_train_nn, X_test_nn, y_train_nn, y_test_nn = train_test(clust_ica, Y_BC)
ica_est_lbl = MLPClassifier(hidden_layer_sizes=(10,5,5), solver='adam',
    ↪activation='relu', learning_rate_init=0.01, max_iter=10000, random_state=18)
train_samp_ica_lbl, NN_train_score_ica_lbl, NN_fit_time_ica_lbl,
    ↪NN_pred_time_ica_lbl, NN_cv_score_ica_lbl = evaluate_nn(ica_est_lbl,
    ↪X_train_nn, y_train_nn, title="Neural Net Wine Quality: ICA")
classification_report(ica_est_lbl, X_train_nn, X_test_nn, y_train_nn, y_test_nn,
    ↪)
```





Classification Report

Training Time: 1.36675



Prediction Time: 0.00084

F1 Score: 0.77

Accuracy: 0.67 AUC: 0.59

Precision: 0.69 Recall: 0.87


| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.58 | 0.32 | 0.41 | 470 |
| 1 | 0.69 | 0.87 | 0.77 | 830 |
| accuracy | | | 0.67 | 1300 |
| macro avg | 0.63 | 0.59 | 0.59 | 1300 |
| weighted avg | 0.65 | 0.67 | 0.64 | 1300 |




Randomized Projection

```
[50]: X_train_nn, X_test_nn, y_train_nn, y_test_nn = train_test(clust_RP, Y_BC)
RP_est_lbl = MLPClassifier(hidden_layer_sizes=(10,5,5), solver='adam',
    ↳activation='relu', learning_rate_init=0.01, max_iter=10000, random_state=18)
train_samp_RP_lbl, NN_train_score_RP_lbl, NN_fit_time_RP_lbl,
    ↳NN_pred_time_RP_lbl, NN_cv_score_RP_lbl = evaluate_nn(RP_est_lbl,
    ↳X_train_nn, y_train_nn, title="Neural Net Wine Quality: Randomize Projection")
classification_report(RP_est_lbl, X_train_nn, X_test_nn, y_train_nn, y_test_nn )
```

Learning Curve For Nueral Network: Neural Net Wine Quality: Randomize Projection



Neural Network model time: Neural Net Wine Quality: Randomize Projection



Classification Report

Training Time: 1.51111


Prediction Time: 0.00079


F1 Score: 0.74

Accuracy: 0.67 AUC: 0.63

Precision: 0.72 Recall: 0.77

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.56 | 0.50 | 0.53 | 483 |
| 1 | 0.72 | 0.77 | 0.74 | 817 |
| accuracy | | | 0.67 | 1300 |
| macro avg | 0.64 | 0.63 | 0.64 | 1300 |
| weighted avg | 0.66 | 0.67 | 0.66 | 1300 |







Random Forest Classifier

```
[51]: X_train_nn, X_test_nn, y_train_nn, y_test_nn = train_test(clust_rfc, Y_BC)
rfc_est_lbl = MLPClassifier(hidden_layer_sizes=(10,5,5), solver='adam',
    ↪activation='logistic', learning_rate_init=0.01,
    ↪max_iter=10000,random_state=18)
train_samp_rfc_lbl, NN_train_score_rfc_lbl, NN_fit_time_rfc_lbl,
    ↪NN_pred_time_rfc_lbl, NN_cv_score_rfc_lbl = evaluate_nn(rfc_est_lbl,
    ↪X_train_nn, y_train_nn,title="Neural Net Wine Quality: Random Forest
    ↪Classifier")
classification_report(rfc_est_lbl, X_train_nn, X_test_nn, y_train_nn, y_test_nn)
```

Learning Curve For Nueral Network: Neural Net Wine Quality: Random Forest Classifier



Neural Network model time: Neural Net Wine Quality: Random Forest Classifier



Classification Report

Training Time: 2.29691


Prediction Time: 0.00053


F1 Score: 0.77

Accuracy: 0.69 AUC: 0.65

Precision: 0.72 Recall: 0.81


| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.61 | 0.49 | 0.54 | 491 |
| 1 | 0.72 | 0.81 | 0.77 | 809 |
| accuracy | | | 0.69 | 1300 |
| macro avg | 0.67 | 0.65 | 0.66 | 1300 |
| weighted avg | 0.68 | 0.69 | 0.68 | 1300 |






Comparison plots


```
[52]: plot_fit_time(train_samp_full_lbl, NN_fit_time_full_lbl, NN_fit_time_pca_lbl,↵
↵NN_fit_time_ica_lbl, NN_fit_time_RP_lbl, NN_fit_time_rfc_lbl, 'Wine Quality_↵
↵Dataset')
```




```
[53]: plot_pred_time(train_samp_full_lbl, NN_pred_time_full_lbl,
    ↪ NN_pred_time_pca_lbl, NN_pred_time_ica_lbl, NN_pred_time_RP_lbl,
    ↪ NN_pred_time_rfc_lbl, 'Wine Quality Dataset')
```



```
[54]: plot_learn_time(train_samp_full_lbl, NN_train_score_full_lbl,␣
      ↪ NN_train_score_pca_lbl, NN_train_score_ica_lbl, NN_train_score_RP_lbl,␣
      ↪ NN_train_score_rfc_lbl, 'Wine Quality Dataset')
```

```
[55]: plot_cv_time(train_samp_full_lbl, NN_cv_score_full_lbl, NN_cv_score_pca_lbl,
↳ NN_cv_score_ica_lbl, NN_cv_score_RP_lbl, NN_cv_score_rfc_lbl, 'Wine Quality_
↳ Dataset')
```



[]:

[]:

[]:

[]: